



Exploring QoS in MapReduce Task Scheduling

Pedro Alvarez-Tabio
Jesus Hernandez
Javier Villa
CS550



Index

- Background
- Motivation
- Overview of MapReduce task scheduling approaches
- Our approach: QoS
- System architecture
- QoS specifications and Service Levels
- Preliminary results
- Conclusions



Background

- MapReduce: data-driven programming model proposed by Google in 2004
 - Apache Hadoop: open-source version
- Especially well suited for distributed data analytics applications.
- Numerous companies and institutions focused on Big Data are relying in these technologies
 - To boost their data processing and computation performance in such data intensive jobs
 - Small tasks can be processed by the nodes in a parallel fashion
 - Example: LinkedIn uses Hadoop MapReduce for discovering People You May Know



Motivation

- Many companies require heavy data processing with MapReduce
- Outsourcing of MapReduce computation to other companies
 - Example: Cloudera
- Current task scheduling techniques in Hadoop do not take into account service guarantees



Overview of Hadoop task scheduling approaches

○ FIFO scheduler

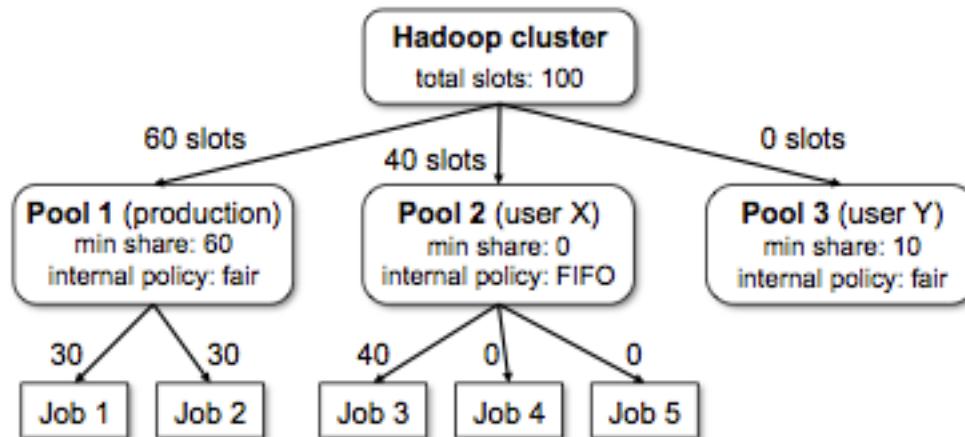
- The original, primitive implementation of task scheduling
- Included in Hadoop by default
- Five priority levels, heartbeat-guided
- Not fair because of its FIFO nature: short, high-priority jobs get stuck behind long, low-priority ones which have been scheduled before



Overview of Hadoop task scheduling approaches (II)

● Fair Scheduler

- Takes into account both fairness and data locality to multiplex clusters efficiently
- Two-level scheduling hierarchy
 - Top level: allocation of tasks across pools using weighted fair sharing
 - Lower level: each pool allocating slots (FIFO or fair sharing)
- Prioritizes meeting minimum shares over fair shares





Overview of Hadoop task scheduling approaches (III)

○ Capacity Scheduler

- Based on arranging jobs into different queues
- Each of them has an allocated fraction of the cluster capacity
- The ordering of jobs can be performed either by a FIFO policy or by prioritizing jobs



Our approach: QoS

- Current approaches aim to specify fairness among jobs
- Neither of them say anything about guaranteed job/task processing
- Our approach: Quality of Service specification in Job and Task scheduler
 - Different client profiles for different performance requirements
 - Permits the fine tuning of the performance in each client
 - Allows the service provider to easily calculate the needed scale of the system in advance

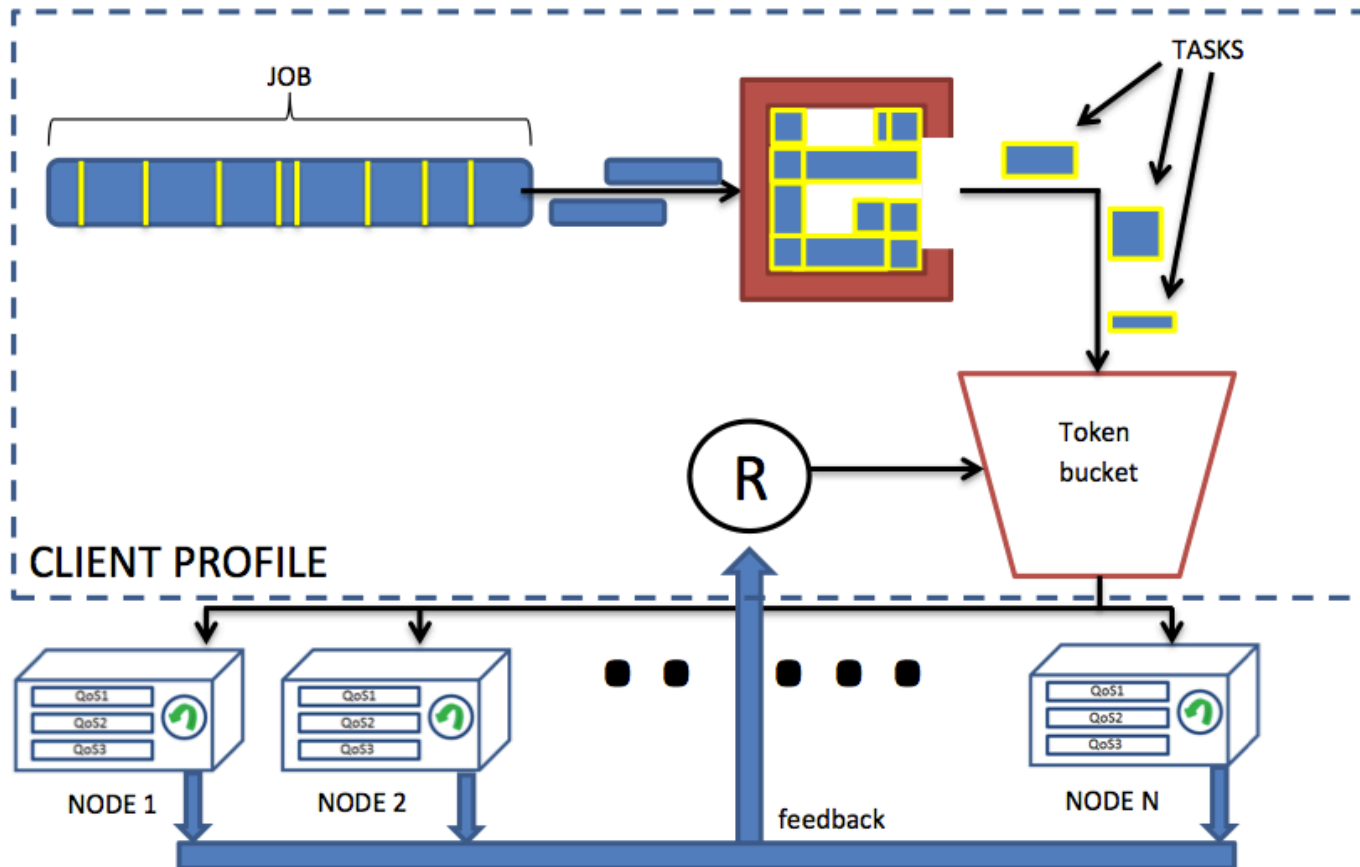


Our approach: main concepts

- We take algorithms from traffic control in computer networks and apply them to the special case of task scheduling
 - Token Bucket
 - Algorithm used in computer networks
 - Checks if data transmissions are conformant to limits in bandwidth and burstiness
 - Weighted Fair Queueing
 - Packet scheduling technique which allows different scheduling priorities (weights) to statistically multiplexed data flows
 - Can be implemented as a weighted round robin



System architecture



QoS Task scheduler



System architecture

◉ Queued Token Bucket

- Tasks are queued before entering the bucket
- Constant token arrival rate
- Non-conformant tasks are sent lowest-priority

◉ Adaptive feedback to clients

- 1 token = 1 ms of processing in any node
- Weighted mean of the last durations gives N tokens needed for next task submission
- This way, we avoid greedy users submitting longer tasks take advantage of the system against other users



System architecture

○ Weighted Fair Queueing

- Each node has a queue for each user profile, with a specified priority
- Weighted round-robin approach
 - Priorities are weights
 - Example: a queue with priority 10 will be polled 10 times a queue with priority 1, in case of great occupation
 - In case one of the queues is empty, the rest share the node's capacity
 - Non-conformant tasks get degraded to a lower-priority queue



QoS specification and Service Classes

- We have defined three Service Classes for the simulator:
 - Premium
 - Token Bucket: 1000 tokens/sec, 10000 tokens max
 - Priority: 10
 - Advanced
 - Token Bucket: 800 tokens/sec, 300 tokens max
 - Priority: 3
 - Basic
 - Token Bucket: 10 token/sec, 100 tokens max
 - Priority: 1



QoS metrics in simulator

- We have defined some metrics to evaluate the system's performance
 - Average job completion time
 - Task processing delay for each Class of Service
 - QoS : performance ratios
 - Percentage of non-conformant tasks



Conclusions

- Our design represents an initial attempt to make MapReduce “as a service” more flexible, scalable and cheap, both for the user and the service provider
- Final steps
 - Complete the metrics and curate the results of those metrics in our simulator



Questions?



THANK YOU